



# **Recapitulare comanda SQL SELECT .**

## **Aspecte avansate**

**- 3 -**

***Interogari combinate. Operatii pe multimi.***

# Tables Used in the Course

## EMPLOYEES

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	SALARY	COMMISSION_PCT	DEPARTMENT_ID	EMAIL	PHONE_NUMBER	HIRE_DATE
100	Steven	King	24000	(null)	90	SKING	515.123.4567	17-JUN-87
101	Neena	Kochhar	17000	(null)	90	NKOCHHAR	515.123.4568	21-SEP-89
102	Lex	De Haan	17000	(null)	90	LDEHAAN	515.123.4569	13-JAN-93
103	Alexander	Hunold	9000	(null)	60	AHUNOLD	590.423.4567	03-JAN-90
104	Bruce	Ernst	6000	(null)	60	BERNST	590.423.4568	21-MAY-91
107	Diana	Lorentz	4200	(null)	60	DLORENTZ	590.423.5567	07-FEB-99
124	Kevin	Mourgos	5800	(null)	50	KMOURGOS	650.123.5234	16-NOV-99
141	Trenna	Rajs	3500	(null)	50	TRAJS	650.121.8009	17-OCT-95
142	Curtis	Davies	3100	(null)	50	CDAVIES	650.121.2994	29-JAN-97
143	Randall	Matos	2600	(null)	50	RMATOS	650.121.2874	15-MAR-98
144	Peter	Vargas	2500	(null)	50	PVARGAS	650.121.2004	09-JUL-98
149	Eleni	Zlotkey	10500	0.2	80	EZLOTKEY	011.44.1344.429018	29-JAN-00
174	Ellen	Abel	11000	0.3	80	EABEL	011.44.1644.429267	11-MAY-96
176	Jonathon	Taylor	8600	0.2	80	JTAYLOR	011.44.1644.429265	24-MAR-98
178	Kimberely	Grant	7000	0.15	(null)	KGRANT	011.44.1644.429263	24-MAY-99
200	Jennifer	Whalen	4400	(null)	10	JWHALEN	515.123.4444	17-SEP-87
201	Michael	Hartstein	13000	(null)	20	MHARTSTE	515.123.5555	17-FEB-96
202	Pat	Fay	6000	(null)	20	PFAY	603.123.6666	17-AUG-97
205	Shelley	Higgins	12000	(null)	110	SHIGGINS	515.123.8080	07-JUN-94
206	William	Gietz	8300	(null)	110	WGIETZ	515.123.8181	07-JUN-94

GRADE_LEVEL	LOWEST_SAL	HIGHEST_SAL
A	1000	2999
B	3000	5999
C	6000	9999
D	10000	14999
E	15000	24999
F	25000	40000

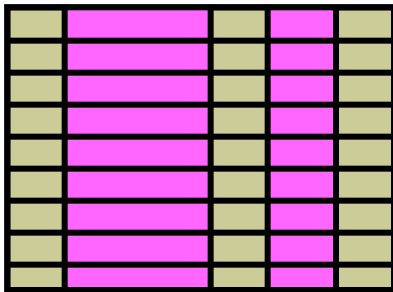
## JOB\_GRADES

DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
10	Administration	200	1700
20	Marketing	201	1800
50	Shipping	124	1500
60	IT	103	1400
80	Sales	149	2500
90	Executive	100	1700
110	Accounting	205	1700
190	Contracting	(null)	1700

## DEPARTMENTS

# Capabilities of SQL `SELECT` Statements

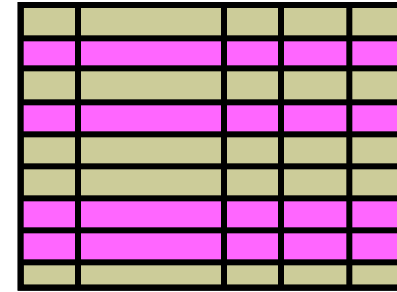
Projection



A 10x5 grid representing a table. The second, third, and fourth columns are highlighted in pink, illustrating the result of a projection operation that selects specific columns from the original table.

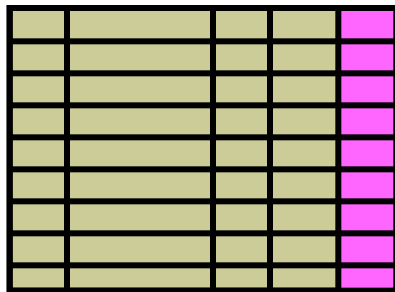

Table 1

Selection



A 10x5 grid representing a table. The first, third, and fourth rows are highlighted in pink, illustrating the result of a selection operation that filters specific rows from the original table.

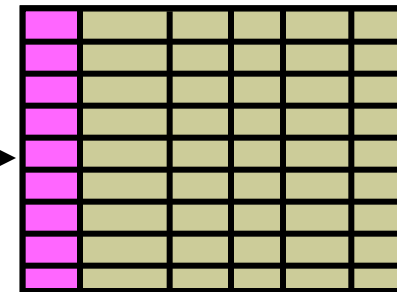

Table 1



A 10x5 grid representing a table. The last column is highlighted in pink, representing a table used in a join operation.


Table 1



Join




A 10x5 grid representing a table. The first column is highlighted in pink, representing a table used in a join operation.


Table 2


# Set Operators



UNION/UNION ALL



INTERSECT



MINUS



# Set Operator Guidelines


- ❑ The expressions in the `SELECT` lists must match in number.
- ❑ The data type of each column in the second query must match the data type of its corresponding column in the first query.
- ❑ Parentheses can be used to alter the sequence of execution (by default, all operators have equal precedence)
- ❑ `ORDER BY` clause can appear only at the very end of the statement.



# Oracle Server and Set Operators

- ❑ Duplicate rows are automatically eliminated except in `UNION ALL`.
- ❑ Column names from the first query appear in the result.
- ❑ The output is sorted in ascending order by default except in `UNION ALL`.

# UNION Operator



The `UNION` operator returns rows from both queries after eliminating duplications.

# Using the UNION Operator

- *Display the current and previous job details of all employees. Display each employee only once.*

```
SELECT employee_id, job_id
FROM   employees
UNION
SELECT employee_id, job_id
FROM   job_history;
```

	EMPLOYEE_ID	JOB_ID
1	100	AD_PRES
2	101	AC_ACCOUNT

...


22	200	AC_ACCOUNT
23	200	AD_ASST

...

27	205	AC_MGR
28	206	AC_ACCOUNT



# UNION ALL Operator



The `UNION ALL` operator returns rows from both queries, including all duplications.


# Using the UNION ALL Operator

- *Display the current and previous departments of all employees.*

```
SELECT employee_id, job_id, department_id
FROM   employees
UNION ALL
SELECT employee_id, job_id, department_id
FROM   job_history
ORDER BY employee_id;
```

	EMPLOYEE_ID	JOB_ID	DEPARTMENT_ID
1	100	AD_PRES	90
...			
17	149	SA_MAN	80
18	174	SA_REP	80
19	176	SA_REP	80
20	176	SA_MAN	80
21	176	SA_REP	80
22	178	SA_REP	(null)
23	200	AD_ASST	10
...			
30	206	AC_ACCOUNT	110

# INTERSECT **Operator**



The `INTERSECT` operator returns rows that are common to both queries.

# Using the INTERSECT Operator

- *Display the employee IDs and job IDs of those employees who currently have a job title that is the same as their previous one (that is, they changed jobs but have now gone back to doing the same job they did previously).*

```
SELECT employee_id, job_id
FROM   employees
INTERSECT
SELECT employee_id, job_id
FROM   job_history;
```

	 EMPLOYEE_ID	 JOB_ID
1	176	SA_REP
2	200	AD_ASST

# Matching the SELECT Statements

- ❑ Using the UNION operator, display the location ID, department name, and the state where it is located.
- ❑ You must match the data type (using the TO\_CHAR function or any other conversion functions) when columns do not exist in one or the other table.

```
SELECT location_id, department_name "Department",  
       TO_CHAR(NULL) "Warehouse location"  
FROM departments  
UNION  
SELECT location_id, TO_CHAR(NULL) "Department",  
       state_province  
FROM locations;
```

# Matching the SELECT Statement

- *Using the UNION operator, display the employee ID, job ID, and salary of all employees.*

```
SELECT employee_id, job_id, salary
FROM   employees
UNION
SELECT employee_id, job_id, 0
FROM   job_history;
```

	EMPLOYEE_ID	JOB_ID	SALARY
1	100	AD_PRES	24000
2	101	AC_ACCOUNT	0
3	101	AC_MGR	0
4	101	AD_VP	17000
5	102	AD_VP	17000
...			
29	205	AC_MGR	12000
30	206	AC_ACCOUNT	8300



# Using the `ORDER BY` Clause in Set Operations

- ❑ The `ORDER BY` clause can appear only once at the end of the compound query.
- ❑ Component queries cannot have individual `ORDER BY` clauses.
- ❑ The `ORDER BY` clause recognizes only the columns of the first `SELECT` query.
- ❑ By default, the first column of the first `SELECT` query is used to sort the output in an ascending order.



# **Recapitulare comanda SQL SELECT .**





## **Aspecte avansate**

**- 4 -**

***lerarhii***



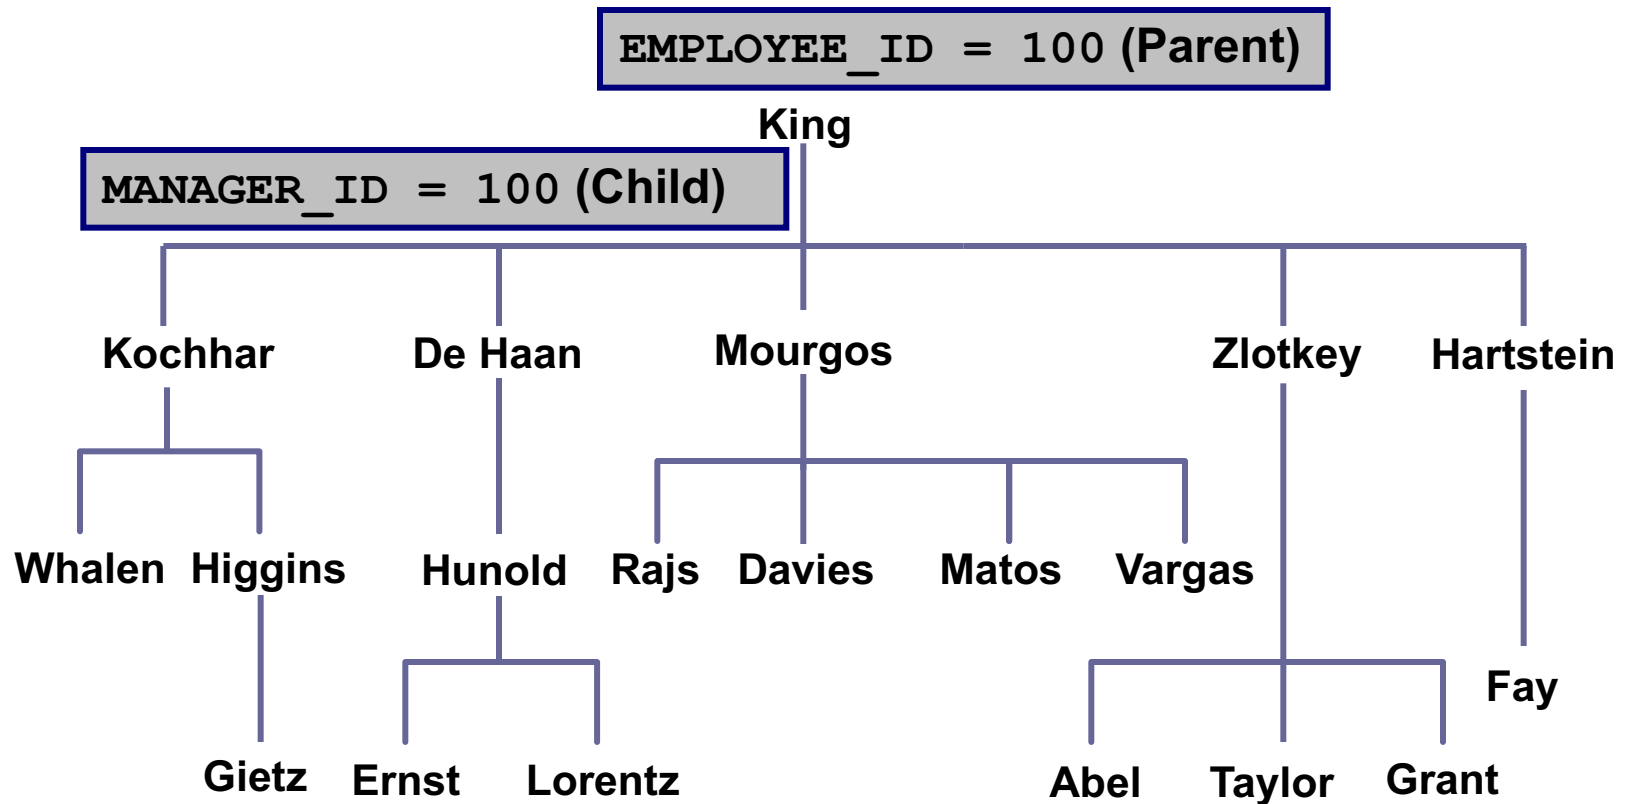
# Sample Data from the EMPLOYEES Table

	 EMPLOYEE_ID	 LAST_NAME	 JOB_ID	 MANAGER_ID
1	100	King	AD_PRES	(null)
2	101	Kochhar	AD_VP	100
3	102	De Haan	AD_VP	100
4	103	Hunold	IT_PROG	102
5	104	Ernst	IT_PROG	103
6	107	Lorentz	IT_PROG	103

...

16	200	Whalen	AD_ASST	101
17	201	Hartstein	MK_MAN	100
18	202	Fay	MK_REP	201
19	205	Higgins	AC_MGR	101
20	206	Gietz	AC_ACCOUNT	205

# Natural Tree Structure



# Hierarchical Queries

```
SELECT [LEVEL], column, expr...  
FROM   table  
[WHERE condition(s)]  
[START WITH condition(s)]  
[CONNECT BY PRIOR condition(s)] ;
```

*condition:*

```
expr comparison_operator expr
```

# Walking the Tree

## Starting Point

- ☐ Specifies the condition that must be met
- ☐ Accepts any valid condition

```
START WITH column1 = value
```

- *Using the EMPLOYEES table, start with the employee whose last name is Kochhar.*

```
...START WITH last_name = 'Kochhar'
```

# Walking the Tree

```
CONNECT BY PRIOR column1 = column2
```

```
... CONNECT BY PRIOR employee_id = manager_id
```

## Direction





Top down      →      Column1 = Parent Key  
                                 Column2 = Child Key

Bottom up      →      Column1 = Child Key  
                                 Column2 = Parent Key

# Walking the Tree: From the Bottom Up

- *Walk from the bottom up, using the `EMPLOYEES` table.*

```
SELECT employee_id, last_name, job_id, manager_id
FROM   employees
START WITH employee_id = 101
CONNECT BY PRIOR manager_id = employee_id ;
```

	 EMPLOYEE_ID	 LAST_NAME	 JOB_ID	 MANAGER_ID
1	101	Kochhar	AD_VP	100
2	100	King	AD_PRES	(null)

# Walking the Tree: From the Top Down


```
SELECT last_name||' reports to '||  
PRIOR last_name "Walk Top Down"  
FROM employees  
START WITH last_name = 'King'  
CONNECT BY PRIOR employee_id = manager_id ;
```

	Walk Top Down
1	King reports to
2	King reports to
3	Kochhar reports to King
4	Greenberg reports to Kochhar
5	Faviet reports to Greenberg

...

105	Grant reports to Zlotkey
106	Johnson reports to Zlotkey
107	Hartstein reports to King
108	Fay reports to Hartstein

# Ranking Rows with the LEVEL Pseudocolumn



Level 1  
root/  
parent

Level 2  
parent/  
child/leaf

Level 3  
parent/  
child/leaf

Level 4  
leaf



# Formatting Hierarchical Reports by Using LEVEL and LPAD


- *Create a report displaying company management levels, beginning with the highest level and indenting each of the following levels.*

```
COLUMN org_chart FORMAT A12
SELECT LPAD(last_name, LENGTH(last_name)+(LEVEL*2)-2, '_')
        AS org_chart
FROM   employees
START WITH first_name='Steven' AND last_name='King'
CONNECT BY PRIOR employee_id=manager_id
```

# Pruning Branches


Use the `WHERE` clause  
to eliminate a node.

```
WHERE last_name != 'Higgins'
```



Use the `CONNECT BY` clause  
to eliminate a branch.

```
CONNECT BY PRIOR  
employee_id = manager_id  
AND last_name != 'Higgins'
```





# **Recapitulare comanda SQL SELECT .**

## **Aspecte avansate**

### **- 5 -**

***Gruparea datelor in rapoarte***



# Reporting operations

- ROLLUP operation to produce subtotal values
- CUBE operation to produce cross-tabulation values
- GROUPING function to identify the row values created by ROLLUP or CUBE
- GROUPING SETS to produce a single result set

# Group Functions: Review

- Group functions operate on sets of rows to give one result per group.

```
SELECT      [column,] group_function(column) . . .
FROM        table
[WHERE      condition]
[GROUP BY  group_by_expression]
[ORDER BY   column];
```

- Example:

```
SELECT AVG(salary), STDDEV(salary),
       COUNT(commission_pct), MAX(hire_date)
FROM   employees
WHERE  job_id LIKE 'SA%';
```

# GROUP BY Clause: Review

## □ Syntax:

```
SELECT      [column,] group_function(column) . . .  
FROM        table  
[WHERE      condition]  
[GROUP BY   group_by_expression]  
[ORDER BY   column];
```

## □ Example:

```
SELECT      department_id, job_id, SUM(salary),  
            COUNT(employee_id)  
FROM        employees  
GROUP BY    department_id, job_id ;
```

# HAVING Clause: Review

- Use the `HAVING` clause to specify which groups are to be displayed.
- You further restrict the groups on the basis of a limiting condition.

```
SELECT      [column,] group_function(column) ...  
FROM        table  
[WHERE      condition]  
[GROUP BY   group_by_expression]  
[HAVING     having_expression]  
[ORDER BY   column];
```



## GROUP BY **with** ROLLUP and CUBE **Operators**

- Use ROLLUP or CUBE with GROUP BY to produce superaggregate rows by cross-referencing columns.
- ROLLUP grouping produces a result set containing the regular grouped rows and the subtotal values.
- CUBE grouping produces a result set containing the rows from ROLLUP and cross-tabulation rows.



# ROLLUP Operator

- ❑ ROLLUP is an extension to the GROUP BY clause.
- ❑ Use the ROLLUP operation to produce cumulative aggregates, such as subtotals.

```
SELECT      [column,] group_function(column) . . .  
FROM        table  
[WHERE      condition]  
[GROUP BY   [ROLLUP] group_by_expression]  
[HAVING     having_expression];  
[ORDER BY   column];
```

# ROLLUP Operator: Example

```
SELECT    department_id, job_id, SUM(salary)
FROM      employees
WHERE     department_id < 60
GROUP BY ROLLUP(department_id, job_id);
```

	DEPARTMENT_ID	JOB_ID	SUM(SALARY)
1	10	AD_ASST	4400
2	10	(null)	4400
3	20	MK_MAN	13000
4	20	MK_REP	6000
5	20	(null)	19000
6	30	PU_MAN	11000
7	30	PU_CLERK	13900
8	30	(null)	24900
9	40	HR_REP	6500
10	40	(null)	6500
11	50	ST_MAN	36400
12	50	SH_CLERK	64300
13	50	ST_CLERK	55700
14	50	(null)	156400
15	(null)	(null)	211200

1

2

3

# CUBE Operator

- ❑ CUBE is an extension to the GROUP BY clause.
- ❑ You can use the CUBE operator to produce cross-tabulation values with a single SELECT statement.

```
SELECT      [column,] group_function(column) ...  
FROM        table  
[WHERE      condition]  
[GROUP BY   [CUBE] group_by_expression]  
[HAVING     having_expression]  
[ORDER BY   column];
```

# CUBE Operator: Example

```
SELECT    department_id, job_id, SUM(salary)
FROM      employees
WHERE     department_id < 60
GROUP BY CUBE (department_id, job_id) ;
```

	DEPARTMENT_ID	JOB_ID	SUM(SALARY)
1	(null)	(null)	211200
2	(null)	HR_REP	6500
3	(null)	MK_MAN	13000
4	(null)	MK_REP	6000
5	(null)	PU_MAN	11000
6	(null)	ST_MAN	36400
7	(null)	AD_ASST	4400
8	(null)	PU_CLERK	13900
9	(null)	SH_CLERK	64300
10	(null)	ST_CLERK	55700
11	10	(null)	4400
12	10	AD_ASST	4400
13	20	(null)	19000
14	20	MK_MAN	13000
15	20	MK_REP	6000
16	30	(null)	24900

1

2

3

4

# GROUPING Function

## ■ The GROUPING function:

- Is used with either the CUBE or ROLLUP operator
- Is used to find the groups forming the subtotal in a row
- Is used to differentiate stored NULL values from NULL values created by ROLLUP or CUBE
- Returns 0 or 1

```
SELECT      [column,] group_function(column) .. ,  
            GROUPING(expr)  
FROM        table  
[WHERE      condition]  
[GROUP BY  [ROLLUP][CUBE] group_by_expression]  
[HAVING    having_expression]  
[ORDER BY  column];
```

# GROUPING Function: Example

```
SELECT    department_id DEPTID, job_id JOB,
          SUM(salary),
          GROUPING(department_id) GRP_DEPT,
          GROUPING(job_id) GRP_JOB
FROM      employees
WHERE     department_id < 50
GROUP BY  ROLLUP(department_id, job_id);
```

	DEPTID	JOB	SUM(SALARY)	GRP_DEPT	GRP_JOB
1	10	AD_ASST	4400	0	0
2	10	(null)	4400	0	1
3	20	MK_MAN	13000	0	0
4	20	MK_REP	6000	0	0
5	20	(null)	19000	0	1
6	30	PU_MAN	11000	0	0
7	30	PU_CLERK	13900	0	0
8	30	(null)	24900	0	1
9	40	HR_REP	6500	0	0
10	40	(null)	6500	0	1
11	(null)	(null)	54800	1	1



# GROUPING SETS

- The `GROUPING SETS` syntax is used to define multiple groupings in the same query.
- All groupings specified in the `GROUPING SETS` clause are computed and the results of individual groupings are combined with a `UNION ALL` operation.
- Grouping set efficiency:
  - Only one pass over the base table is required.
  - There is no need to write complex `UNION` statements.
  - The more elements `GROUPING SETS` has, the greater is the performance benefit.

# GROUPING SETS: Example

```
SELECT    department_id, job_id,  
          manager_id,AVG(salary)  
FROM      employees  
GROUP BY GROUPING SETS  
         ((department_id,job_id), (job_id,manager_id));
```

R2	DEPARTMENT_ID	R2	JOB_ID	R2	MANAGER_ID	R2	AVG(SALARY)
1		(null)	SH_CLERK		122		320
2		(null)	AC_MGR		101		1200
3		(null)	ST_MAN		100		728
4	...	(null)	ST_CLERK		121		267

1

R2	DEPARTMENT_ID	R2	JOB_ID	R2	MANAGER_ID	R2	AVG(SALARY)
39	110	AC_MGR		(null)		12000	
40	90	AD_PRES		(null)		24000	
41	60	IT_PROG		(null)		5760	
42	100	FI_MGR		(null)		12000	

2

...



# Composite Columns

- A composite column is a collection of columns that are treated as a unit.

```
ROLLUP (a, (b, c) , d)
```

- Use parentheses within the GROUP BY clause to group columns, so that they are treated as a unit while computing ROLLUP or CUBE operations.
- When used with ROLLUP or CUBE, composite columns would require skipping aggregation across certain levels.

# Composite Columns: Example

```
SELECT    department_id, job_id, manager_id,  
          SUM(salary)  
FROM      employees  
GROUP BY ROLLUP( department_id, (job_id, manager_id)) ;
```

1	DEPARTMENT_ID	JOB_ID	MANAGER_ID	SUM(SALARY)	2
	(null)	SA_REP	149	7000	
	(null)	(null)	(null)	7000	
	10	AD_ASST	101	4400	
	10	(null)	(null)	4400	
	20	MK_MAN	100	13000	
	20	MK_REP	201	6000	
	20	(null)	(null)	19000	
...					
	100	FI_MGR	101	12000	
	100	FI_ACCOUNT	108	39600	
	100	(null)	(null)	51600	3
	110	AC_MGR	101	12000	
	110	AC_ACCOUNT	205	8300	
	110	(null)	(null)	20300	
	(null)	(null)	(null)	691400	4

# Concatenated Groupings

- Concatenated groupings offer a concise way to generate useful combinations of groupings.
- To specify concatenated grouping sets, you separate multiple grouping sets, `ROLLUP` and `CUBE` operations with commas so that the Oracle Server combines them into a single `GROUP BY` clause.
- The result is a cross-product of groupings from each `GROUPING SET`.

```
GROUP BY GROUPING SETS (a, b) , GROUPING SETS (c, d)
```

# Concatenated Groupings: Example

```
SELECT  department_id, job_id, manager_id,  
        SUM(salary)  
FROM    employees  
GROUP BY department_id,  
        ROLLUP(job_id),  
        CUBE(manager_id);
```

	DEPARTMENT_ID	JOB_ID	MANAGER_ID	SUM(SALARY)
1	(null)	SA_REP	149	7000
2	10	AD_ASST	101	4400
3	20	MK_MAN	100	13000
4	20	MK_REP	201	6000

1

...

90	AD_VP	100	34000
90	AD_PRES	(null)	24000

...

(null)	SA_REP	(null)	7000
10	AD_ASST	(null)	4400

2

...

110	(null)	101	12000
110	(null)	205	8300
110	(null)	(null)	20300

3